

# Networking In the Linux Kernel (kerNet)

Lecture 11

TCP Structures In the Kernel, Assortment of Sockets

# Announcements

---

- Peer 1 Reviews available
  - Click “View Grade” by Peer Review Return: Assignment 1
  - Download your PDFs
- Assignment 2 due Thursday night
  - GFP\_USER is not your friend!
  - Pre-allocate your 2-D array in user-space
  - Static arrays may also do weird things: `char** x`; and then manual allocation usually works out better than `char x[256][256]`.

# Assignment 2 Discussion

---

- Why does the sleep() cause issues?

# Testing Discussion

---

- “Basic arguments” – What does this mean?
- “I ran the standard test” - What does this mean?
- “When I tried with invalid arguments” – What does this mean?
- When you run multiple tests in a row, it’s possible they could be related! But the ordering of tests (and the exact details) were rarely shared
- “Multiple clients” / “Ran clients and killed the server” – when? How many clients? Did some/all of them finish writing?

# Testing Discussion

---

- Screenshots are not required but can be helpful
- Exact commands and the order you run them in are also useful
- “I tried with values of x from 0 to 5” ... is that
  - 0, 3, 5?
  - 0,1,2,3,4,5?
  - ???
- OS, compiler version, etc. can be helpful but I’ll consider “extra”
- If you can’t reproduce a bug, it’s still worth providing as much detail as you can – perhaps the original coder will have insight / will be able to run your tests many times and see the bug happen.

# Testing Discussion

---

- Very few students had enough description to be sure I was reproducing their tests correctly (especially when there were multiple runs required).
- I'll be much pickier in future peer reviews about your explanation of testing!
- Communication intensive course – and being able to communicate precisely about testing is important (arguably more so than essay-style writing will be in most of your careers)

# Sources for Today

---

- Kurose, James F. and Ross, W. Keith. Computer Networking A Top-Down Approach. Pearson Education Inc., 2006.
  - Basically just for the TCP header

# Quick Teaser About Windows...

---

*/\* Update our send window.*

*\**

*\* Window update algorithm, described in RFC793/RFC1122 (used in linux-2.2*

*\* and in FreeBSD. NetBSD's one is even worse.) is wrong.*

*\*/*

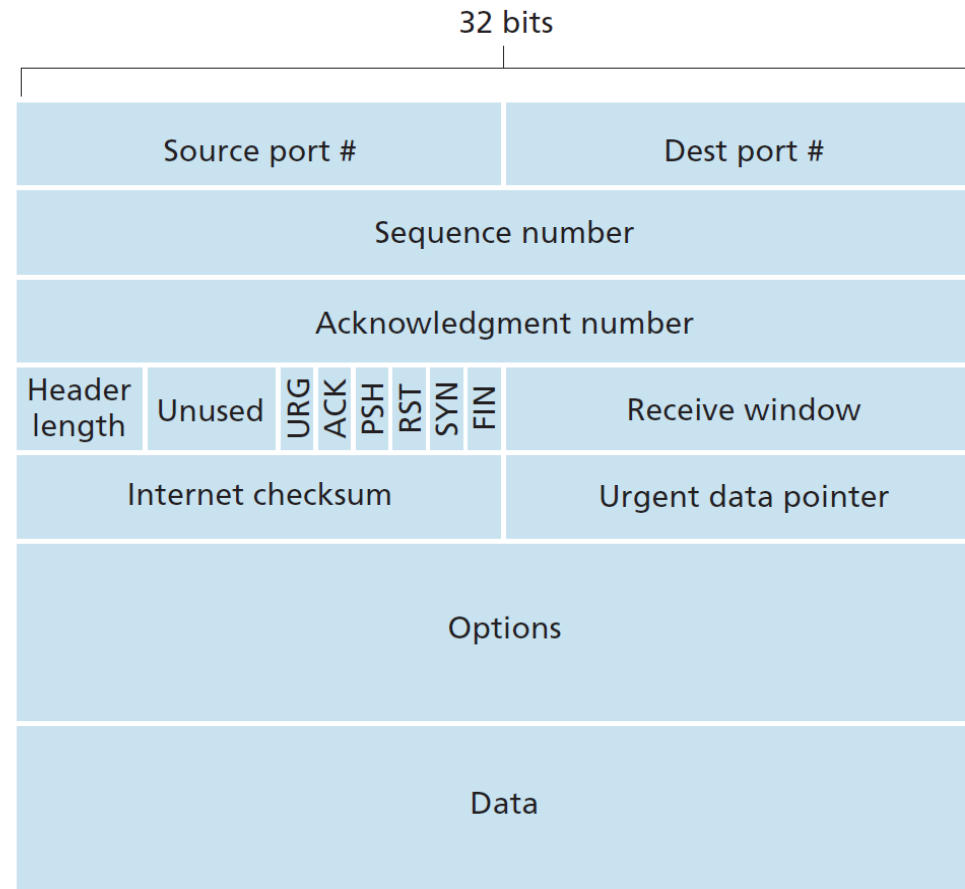
*static int tcp\_ack\_update\_window(struct sock \*sk, const struct sk\_buff \*skb, u32  
ack, u32 ack\_seq)*



# TCP Header

# TCP Segment Structure

---



**Figure 3.29** ♦ TCP segment structure

# A Few Header Notes

---

- This gives us the basic header layout
- We haven't discussed options
- We might have a variable number of TCP options.
  - Not all options have the same length
- Header length is 4 bits, representing the number of 4-byte words the header-with-options takes
  - Without options, we're looking at 20 bytes (5 words), you'll often see 0x5
  - Remember this is a half-byte, so it'll get mashed together with "unused" to yield 0x50. With [RFC 3540](#) you might see 0x51, 9<sup>th</sup> flag bit!

# Header in the Kernel

---

- [tcp\\_hdr](#)(skb) returns a pointer to the transport header data in skb
- The return type is [struct tcphdr](#)\* defined in <linux/uapi/tcp.h>
- Let's pull it up and compare it to the TCP Segment Structure graphic from earlier

# tcphdr (1/3)

---

```
struct tcphdr {  
    __be16    source;  
    __be16    dest;  
    __be32    seq;  
    __be32    ack_seq;
```

# tcphdr (2/3)

---

```
#elif defined(__BIG_ENDIAN_BITFIELD)  
    __u16 doff:4,  
        res1:4,  
        cwr:1,  
        ece:1,  
        urg:1,  
        ack:1,  
        psh:1,  
        rst:1,  
        syn:1,  
        fin:1;
```

# tcphdr (3/3)

---

*#endif*

*\_\_be16 window;*

*\_\_sum16 check;*

*\_\_be16 urg\_ptr;*

*};*

# tcphdr enum

---

```
enum {  
    TCP_FLAG_CWR = __constant_cpu_to_be32(0x00800000),  
    TCP_FLAG_ECE = __constant_cpu_to_be32(0x00400000),  
    TCP_FLAG_URG = __constant_cpu_to_be32(0x00200000),  
    TCP_FLAG_ACK = __constant_cpu_to_be32(0x00100000),  
    TCP_FLAG_PSH = __constant_cpu_to_be32(0x00080000),  
    TCP_FLAG_RST = __constant_cpu_to_be32(0x00040000),  
    TCP_FLAG_SYN = __constant_cpu_to_be32(0x00020000),  
    TCP_FLAG_FIN = __constant_cpu_to_be32(0x00010000),  
    TCP_RESERVED_BITS = __constant_cpu_to_be32(0x0F000000),  
    TCP_DATA_OFFSET = __constant_cpu_to_be32(0xF0000000)  
}; \[source\]
```



# What About TCP Options?

---

- A bunch of #defines for option numbers and lengths in <net/tcp.h>
- Look for the definition of [TCPOPT\\_SACK\\_PERM](#) to find the corresponding blocks
- net/ipv4/tcp\_input.c: [tcp\\_parse\\_options\(\)](#)
  - We normally only do this during handshakes (SYN flag set), slowpath parsing
  - First attempts tcp\_fast\_parse\_options() which only expects a timestamp
  - You'll see *doff* used here, get used to it - this is the data offset.
- If you walk further up the call chain, [tcp\\_fast\\_parse\\_options\(\)](#) or [tcp\\_conn\\_request\(\)](#) or [tcp\\_rcv\\_fastopen\\_synack\(\)](#) or [tcp\\_synsent\\_state\\_process\(\)](#)
- Fast path only called through [tcp\\_validate\\_incoming\(\)](#) which is...
  - Getting off track (for now), a lot about PAWS/RST here
  - Called by [tcp\\_rcv\\_established\(\)](#) and [tcp\\_rcv\\_state\\_process\(\)](#)

# TCP Socket Structure

# tcp\_sock (1/25)

---

```
struct tcp\_sock {  
    /* inet_connection_sock has to be the first member of tcp_sock */  
    struct inet\_connection\_sock inet_conn;  
    u16    tcp_header_len;    /* Bytes of tcp header to send */  
    u16    gso_segs;    /* Max number of segs per GSO packet */  
  
    /*  
    * Header prediction flags  
    * 0x5?10 << 16 + snd_wnd in net byte order  
    */  
    __be32    pred_flags;
```

# Detour: inet\_connection\_sock

---

```
struct inet_connection_sock {  
    /* inet_sock has to be the first member! */  
    struct inet_sock icsk_inet;  
    ...  
}
```

```
struct inet_sock {  
    /* sk and pinet6 has to be the first two members of inet_sock */  
    struct sock      sk;  
#if IS_ENABLED(CONFIG_IPV6)  
    struct ipv6_pinfo *pinet6;  
#endif
```

# Detour: sock

---

```
struct sock {  
    /*  
     * Now struct inet_timewait_sock also uses sock_common, so please just  
     * don't add anything before this first member (__sk_common) --acme  
     */  
    struct sock_common    __sk_common;
```

```
/**  
 * struct sock_common - minimal network layer representation of sockets
```

# tcp\_sock -> sock\_common

Modified from [https://en.wikipedia.org/wiki/Matryoshka\\_doll#/media/File:Russian-Matroschka.jpg](https://en.wikipedia.org/wiki/Matryoshka_doll#/media/File:Russian-Matroschka.jpg), CC BY-SA 3.0



# tcp\_sock (2/25)

---

```
/*
 * RFC793 variables by their proper names. This means you can
 * read the code and the spec side by side (and laugh ...)
 * See RFC793 and RFC1122. The RFC writes these in capitals.
 */
u64      bytes_received;    /* RFC4898 tcpEStatsAppHCThruOctetsReceived
 * sum(delta(rcv_nxt)), or how many bytes
 * were acked.
 */
u32      segs_in;    /* RFC4898 tcpEStatsPerfSegsIn
 * total number of segments in.
 */
u32      data_segs_in;    /* RFC4898 tcpEStatsPerfDataSegsIn
 * total number of data segments in.
 */
```

# tcp\_sock (3/25)

---

```
u32  rcv_nxt;    /* What we want to receive next    */
u32  copied_seq; /* Head of yet unread data                        */
u32  rcv_wup;    /* rcv_nxt on last window update sent
*/
u32  snd_nxt;    /* Next sequence we send                        */
u32  segs_out;   /* RFC4898 tcpEStatsPerfSegsOut
                * The total number of segments sent.
                */
```



# tcp\_sock (4/25)

---

<i>u32</i>	<i>data_segs_out;</i>	<i>/* RFC4898 tcpEStatsPerfDataSegsOut * total number of data segments sent. */</i>
<i>u64</i>	<i>bytes_sent;</i>	<i>/* RFC4898 tcpEStatsPerfHCDataOctetsOut * total number of data bytes sent. */</i>
<i>u64</i>	<i>bytes_acked;</i>	<i>/* RFC4898 tcpEStatsAppHCThruOctetsAcked * sum(delta(snd_una)), or how many bytes * were acked. */</i>
<i>u32</i>	<i>dsack_dups;</i>	<i>/* RFC4898 tcpEStatsStackDSACKDups * total number of DSACK blocks received */</i>

# tcp\_sock (5/25)

---

```
    u32    snd_una;      /* First byte we want an ack for */
    u32    snd_sml;      /* Last byte of the most recently transmitted small
packet */
    u32    rcv_tstamp;   /* timestamp of last received ACK (for keepalives) */
    u32    lsndtime;     /* timestamp of last sent data packet (for restart
window) */
    u32    last_oow_ack_time; /* timestamp of last out-of-window ACK */
    u32    compressed_ack_rcv_nxt;
```

- Minshall's modification to Nagle's algorithm: [\[here\]](#)
- TCP SACK Compression added in 4.19.7: [\[discussion\]](#)

# tcp\_sock (6/25)

---

*u32 tsoffset; /\* timestamp offset \*/*

*struct list\_head tsq\_node; /\* anchor in tsq\_tasklet.head list \*/*

*struct list\_head tsorted\_sent\_queue; /\* time-sorted sent but un-SACKed skbs \*/*

*/\* Data for direct copy to user \*/*

*struct ucopy; <----- no longer exists, see [[prequeue discussion](#)]*

# Detour: tsq?

---

```
/* TCP SMALL QUEUES (TSQ)  
*  
* TSQ goal is to keep small amount of skbs per tcp flow in tx queues (qdisc+dev)  
* to reduce RTT and bufferbloat.  
* We do this using a special skb destructor (tcp_wfree).  
*  
* Its important tcp_wfree() can be replaced by sock_wfree() in the event skb  
* needs to be reallocated in a driver.  
* The invariant being skb->truesize subtracted from sk->sk_wmem_alloc  
*  
* Since transmit from skb destructor is forbidden, we use a tasklet  
* to process all sockets that eventually need to send more skbs.  
* We use one tasklet per cpu, with its own queue of sockets.  
*/  
struct tsq_tasklet {
```

# TSQ Supplemental Reading

---

- <https://lwn.net/Articles/507065/>
  - Has a link to the patch proposal, which is at: <https://lwn.net/Articles/506237/>

# tcp\_sock (7/25)

---

```
u32    snd_wl1;    /* Sequence for window update      */
u32    snd_wnd;    /* The window we expect to receive      */
u32    max_window; /* Maximal window ever seen from peer  */
/*
u32    mss_cache; /* Cached effective mss, not including SACKS */

u32    window_clamp; /* Maximal window to advertise
/*
u32    rcv_ssthresh; /* Current window clamp      */
```

# SND.WL2?

---

- From RFC793:
  - SND.WL1 - segment sequence number used for last window update
  - SND.WL2 - segment acknowledgment number used for last window update
  - These plus some other variables should be in a “Transmission Control Block, or TCB” but we have them in *tcp\_sock*
  - Not to be confused with *tcp\_skb\_cb* (discussed next time)
- Claim: WL2 is redundant
  - Reading: <https://www.ietf.org/mail-archive/web/tsvwg/current/msg03445.html>

# tcp\_sock (8/25)

---

```
/* Information of the most recently (s)acked skb */
struct tcp_rack {
    struct skb_mstamp mstamp; /* (Re)sent time of the skb */
    u8 advanced; /* mstamp advanced since last lost marking */
    u8 reord; /* reordering detected */
} rack;
```

See also, comment above [tcp\\_rack\\_detect\\_loss\(\)](#) and [draft-ietf-tcpm-rack-01](#)

Lots of work on this – two years ago we were on revision 05, last year revision 11.

Current version is a full RFC: [RFC 8985](#)



# tcp\_sock (9/25)

---

```
u16    advmss;                /* Advertised MSS */
u8      compressed_ack;
u8      dup_ack_counter:2,
        tlp_retrans:1, /* TLP is a retransmission */
        unused:5;
u32     chrono_start; /* Start time in jiffies of a TCP chrono */
u32     chrono_stat[3]; /* Time in jiffies for chrono_stat stats */
u8      chrono_type:2,      /* current chronograph type */
        rate_app_limited:1, /* rate_{delivered,interval_us} limited? */
        ....
```

# tcp\_sock (10/25)

---

```
u8    nonagle    : 4,/* Disable Nagle algorithm?      */
      thin_lto    : 1,/* Use linear timeouts for thin streams */
      thin_dupack : 1,/* Fast retransmit on first dupack    */
      repair      : 1,
      frto        : 1;/* F-RTO (RFC5682) activated in CA_Loss */
```

This part actually comes AFTER the next slide, just putting it here for space reasons:

```
u32    tcp_tx_delay; /* delay (in usec) added to TX packets */
u64    tcp_wstamp_ns; /* departure time for next sent data packet */
u64    tcp_clock_cache; /* cache last tcp_clock_ns() (see
tcp_mstamp_refresh()) */
```

# tcp\_sock (11/25)

---

```
u8    repair_queue;
u8    do_early_retrans:1, /* Enable RFC5827 early-retransmit */
      syn_data:1,      /* SYN includes data */
      syn_fastopen:1,   /* SYN includes Fast Open option */
      syn_fastopen_exp:1, /* SYN includes Fast Open exp. option */
      syn_data_acked:1, /* data in SYN is acked by SYN-ACK */
      save_syn:1,      /* Save headers of SYN packet */
      is_cwnd_limited:1; /* forward progress limited by snd_cwnd? */
u32    tlp_high_seq; /* snd_nxt at the time of TLP retransmit. */
```

# tcp\_sock (12/25)

---

*/\* RTT measurement \*/*

*u32 srtt\_us; /\* smoothed round trip time << 3 in usecs \*/*  
*u32 mdev\_us; /\* medium deviation \*/*  
*u32 mdev\_max\_us; /\* maximal mdev for the last rtt period \*/*  
*u32 rttvar\_us; /\* smoothed mdev\_max \*/*  
*u32 rtt\_seq; /\* sequence number to update rttvar \*/*  
*struct minmax rtt\_min;*

# tcp\_sock (13/25)

---

```
u32  packets_out; /* Packets which are "in flight" */
u32  retrans_out; /* Retransmitted packets out */
u32  max_packets_out; /* max packets_out in last window */
u32  max_packets_seq; /* right edge of max_packets_out flight */

u16  urg_data; /* Saved octet of OOB data and control flags */
u8   ecn_flags; /* ECN status bits. */
u8   keepalive_probes; /* num of allowed keep alive probes */
u32  reordering; /* Packet reordering metric. */
u32  snd_up; /* Urgent pointer */
```

# tcp\_sock (14/25)

---

```
/*
 * Options received (usually on last packet, some only on SYN packets).
 */
    struct tcp\_options\_received rx_opt;

/*
 * Slow start and congestion control (see also Nagle, and Karn & Partridge)
 */
    u32      snd_ssthresh;          /* Slow start size threshold */
    u32      snd_cwnd; /* Sending congestion window */
    u32      snd_cwnd_cnt;         /* Linear increase counter */
    u32      snd_cwnd_clamp; /* Do not allow snd_cwnd to grow above this */
    u32      snd_cwnd_used;
    u32      snd_cwnd_stamp;
```

# tcp\_sock (15/25)

---

```
u32    prior_cwnd;    /* Congestion window at start of Recovery. */
u32    prr_delivered; /* Number of newly delivered packets to
                        * receiver in Recovery. */
u32    prr_out;       /* Total number of pkts sent during Recovery. */
u32    delivered;     /* Total data packets delivered incl. rexmits */
u32    lost;          /* Total data packets lost incl. rexmits */
u32    app_limited;    /* limited until "delivered" reaches this val */
u64    first_tx_mstamp; /* start of window send phase */
u64    delivered_mstamp; /* time we reached "delivered" */
u32    rate_delivered; /* saved rate sample: packets delivered */
u32    rate_interval_us; /* saved rate sample: time elapsed */
```

See [[net/ipv4/tcp\\_rate.c](https://net.ipv4/tcp_rate.c)] for more on rate variables

# tcp\_sock (16/25)

---

```
u32  rcv_wnd;    /* Current receiver window          */
u32  write_seq;  /* Tail(+1) of data held in tcp send buffer */
u32  notsent_lowat; /* TCP_NOTSENT_LOWAT */
u32  pushed_seq; /* Last pushed seq, required to talk to
windows */
u32  lost_out;   /* Lost packets */
u32  sacked_out; /* SACK'd packets */
u32  fackets_out; /* FACK'd packets */
```



# tcp\_sock (17/25)

---

```
struct hrtimer      pacing_timer;  
struct hrtimer      compressed_ack_timer;
```

```
/* from STCP, retrans queue hinting */  
struct sk_buff* lost_skb_hint;  
struct sk_buff *retransmit_skb_hint;
```

```
/* OOO segments go in this rbtree. Socket lock must be held. */  
struct rb_root      out_of_order_queue;  
struct sk_buff*ooo_last_skb; /* cache rb_last(out_of_order_queue) */
```

# tcp\_sock (18/25)

---

```
/* SACKs data, these 2 need to be together (see tcp_options_write) */  
struct tcp_sack_block duplicate_sack[1]; /* D-SACK block */  
struct tcp_sack_block selective_acks[4]; /* The SACKS themselves*/  
  
struct tcp_sack_block recv_sack_cache[4];  
  
struct sk_buff *highest_sack; /* skb just after the highest  
    * skb with SACKed bit set  
    * (validity guaranteed only if  
    * sacked_out > 0)  
*/
```

# tcp\_sock (19/25)

---

*int   lost\_cnt\_hint;*

*u32   prior\_ssthresh; /\* ssthresh saved at recovery start   \*/*

*u32   high\_seq;       /\* snd\_nxt at onset of congestion    \*/*

*u32   retrans\_stamp;       /\* Timestamp of the last retransmit,  
                              \* also used in SYN-SENT to remember stamp of  
                              \* the first SYN. \*/*

*u32   undo\_marker; /\* snd\_una upon a new recovery episode. \*/*

*int   undo\_retrans; /\* number of undoable retransmissions. \*/*

# tcp\_sock (20/25)

---

*u64      bytes\_retrans;    /\* RFC4898 tcpEStatsPerfOctetsRetrans*

*\* Total data bytes retransmitted*

*\*/*

*u32      total\_retrans;    /\* Total retransmits for entire connection \*/*

*u32      urg\_seq; /\* Seq of received urgent pointer \*/*

*unsigned int              keepalive\_time;    /\* time before keep alive takes place \*/*

*unsigned int              keepalive\_intvl; /\* time interval between keep alive probes \*/*

*int                        linger2;*

# tcp\_sock (21/25)

---

*u16 timeout\_rehash; /\* Timeout-triggered rehash attempts \*/*

*u32 rcv\_oobpack; /\* Received out-of-order packets, for tcpinfo \*/*

*/\* Receiver side RTT estimation \*/*

*u32 rcv\_rtt\_last\_tsecr;*

*struct {*

*u32 rtt;*

*u32 seq;*

*u32 time;*

*} rcv\_rtt\_est;*

# tcp\_sock (22/25)

---

*/\* Receiver queue space \*/*

*struct {*

*int space;*

*u32 seq;*

*u32 time;*

*} rcvq\_space;*

# tcp\_sock (23/25)

---

```
/* TCP-specific MTU probe information. */
    struct {
        u32          probe_seq_start;
        u32          probe_seq_end;
    } mtu_probe;
    u32  mtu_info; /* We received an ICMP_FRAG_NEEDED /
ICMPV6_PKT_TOOBIG
                * while socket was owned by user.
                */
```

## tcp\_sock (24/25)

---

```
#ifdef CONFIG_TCP_MD5SIG  
/* TCP AF-Specific parts; only used by MD5 Signature support so far */  
    const struct tcp_sock_af_ops    *af_specific;  
  
/* TCP MD5 Signature Option information */  
    struct tcp_md5sig_info    __rcu *md5sig_info;  
#endif
```



# tcp\_sock (25/25)

---

```
/* TCP fastopen related information */  
    struct tcp_fastopen_request *fastopen_req;  
    /* fastopen_rsk points to request_sock that resulted in this big  
    * socket. Used to retransmit SYNACKs etc.  
    */  
    struct request_sock *fastopen_rsk;  
    u32    *saved_syn;  
};
```

See also: [RFC 7413](#)

# For Next Time

---

- Monday: TCP CBs, Connection and States
  - Assignment 2 Peer Review will also be assigned
- Thursday: TCP Congestion Avoidance

## **Looking forward:**

- Monday (10/25): Projects 1 AND 2 will be released
- Monday 10/25, Thursday 10/28: In-class time for Project 1
- Monday 11/1 – Monday 11/8: Project 1 Presentations (these are in-person) and Project 1 Peer Reviews (presentation feedback)