

# Peer-to-Peer Networking

Network Programming

# Interesting Approaches

- High-level overview
  - Only scratch the surface!
- Too many to cover any in depth
  - High level overview
- I encourage everyone to dig more deeply if anything intrigues you

# Motivation

- Routing resilience
  - If a node breaks, should be able to route around the damage
- Communication still feasible
- Formerly, prioritize connectivity
- Now prioritize content from endpoints!

# Problems

- What if an endpoint goes down?
  - Gmail? Facebook? Twitter?
- What if a state-level agency decides Facebook is bad?
  - Maybe Bhutan bans depression sites?
- What about natural disasters?
  - I need to find my family/friends

# Peer-to-Peer

- All peers are equal
  - Content providers
  - Routing partners
- Pay to play
  - Each host generates workload
  - Each host also contributes resources
- Conceptually, scales well


# Examples

- Napster
  - P2P w/ a centralized server
  - Single point of failure!
- Gnutella
  - Creates an overlay network
  - Floods content requests
  - No single point of failure, more resilient than Napster but scalability issues
- Both have limitations!

# Distributed Hash Tables

- Properties:
  - Decentralized
  - Fault Tolerant
  - Scalable
- Generally use keyspace partitioning scheme
- Each node maintains a (partial) routing table for the overlay network

# Routing in Overlay Network





# Chord

- ids are length  $m$  (so entire space is  $0 \dots 2^m - 1$ )
  - Common value is 16-bits (64k keys/nodes)
- Node  $a$  wants to try to find key  $k$
- Who is  $\text{successor}(k)$ , the node responsible for  $k$ ?
  - If a node  $j$  has ID  $k$ , then  $\text{successor}(k) = j$
  - o.w. node  $j'$  (the first node going clockwise from where  $k$  would be, i.e. closest node with  $\text{ID} > k$ ) is the successor
- If  $\text{successor}(k) \neq a$  and  $\neq a+1$ , ask  $a+1$  to find the next node.
- $O(N)$  search, not great!

# Chord – Finger Tables

- Every node keeps a “finger table” with up to  $m$  entries
- Entry  $i$  for node  $n$  is
  - $\text{successor}((n + 2^{i-1}) \bmod 2^m)$
  - Now we can “jump” further along the circle as needed, instead of having to ask  $n+1, n+2, n+3, \dots$
  - $\text{successor}(k)$
  - $O(\log N)$  search time, much better!
  - Should feel a lot like a binary search
- See [Wikipedia](#) for more details

# Gnutella

- Nodes are called servants
- Connect, then send descriptors (Gnutella protocol concept, not C file descriptor)
  - Ping to discover
  - Reply with one or more pongs to describe availability
  - Query to ask for data
  - QueryHit to reply that you have the data

# Gnutella Routing

- Pongs can only travel back along the path the Ping came from
- QueryHit similarly can only follow Query's path
- Forward Ping/Query to all direct connections except the one you received it on
- Decrement TTL field and increment Hops before forwarding. Do not forward if  $TTL = 0$
- Don't forward if payload descriptor and descriptor ID already seen - prevents replication of messages!

# Gnutella Downloads

- Remember that Gnutella is an overlay
- For downloads we do a direct client-to-client connection
  - Don't route through other servants in the overlay network
- This is still P2P, no dedicated server
  - But if we could stop the initial distributor before copies were shared, we could prevent a file from spreading
  - Easier said than done!

# BitTorrent

- Break files up into *pieces*
- Upon downloading a piece, able to then serve it to another consumer
  - Popularity increases number of hosts for that piece / content
- Cryptographic hashing used to ensure no changes to piece made

# Chapter 2

## Application Layer

### A note on the use of these Powerpoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved



## *Computer Networking: A Top Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP




# Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

## *examples:*

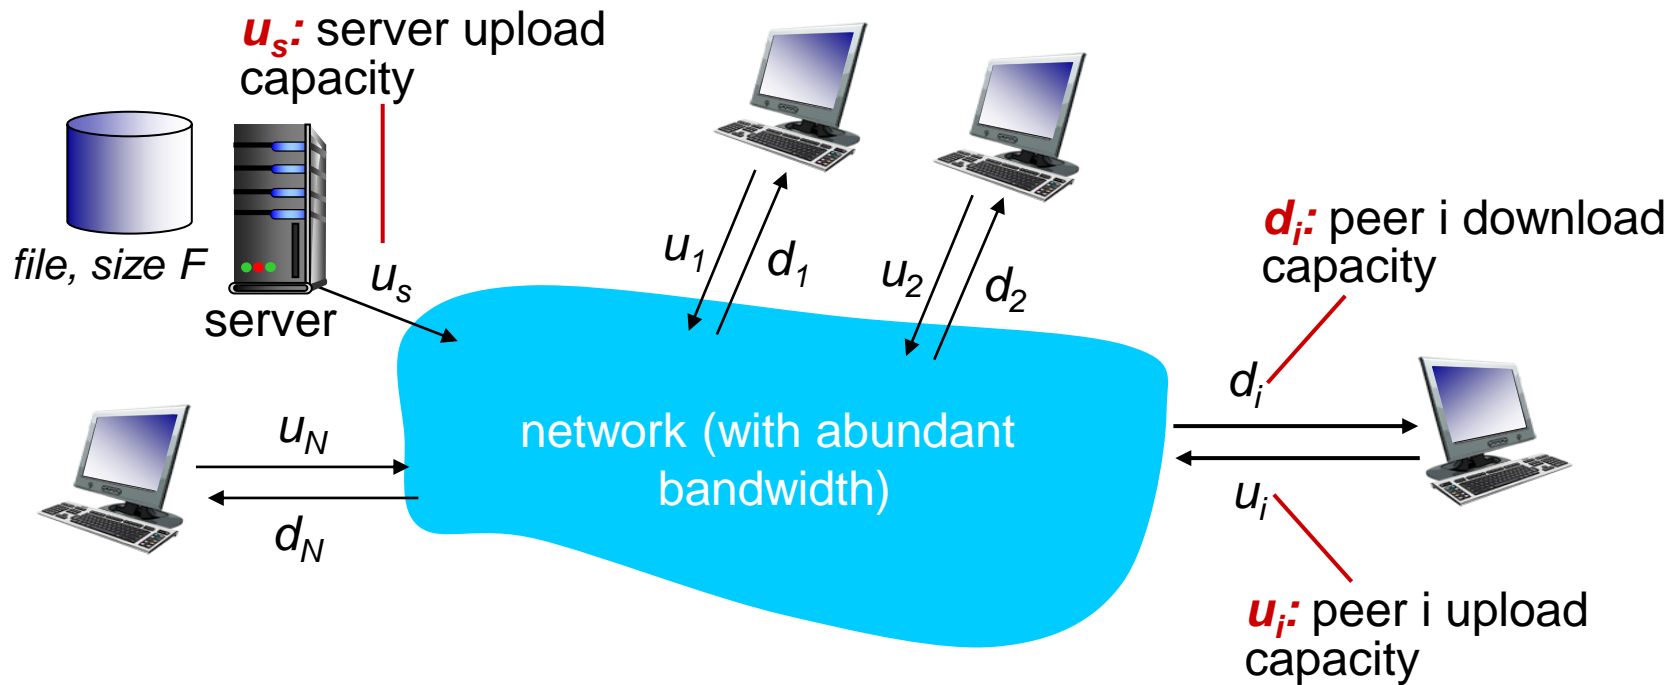
- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



# File distribution: client-server vs P2P

Question: how much time to distribute file (size  $F$ ) from one server to  $N$  peers?

- peer upload/download capacity is limited resource




# Side Question

- If a tracker is required, is the tracker legally responsible?
- What if we killed/blocked the tracker?
  - What effect does killing the tracker have?
  - Does it kill the P2P network?
  - Can we somehow still access content?

# File distribution time: client-server

- **server transmission:** must sequentially send (upload)  $N$  file copies:
  - time to send one copy:  $F/u_s$
  - time to send  $N$  copies:  $NF/u_s$
- **client:** each client must download file copy
  - $d_{min}$  = min client download rate
  - min client download time:  $F/d_{min}$




*time to distribute  $F$   
to  $N$  clients using  
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$

# File distribution time: P2P

- **server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- **client:** each client must download file copy
  - min client download time:  $F/d_{\min}$
- **clients:** as aggregate must download  $NF$  bits
  - max upload rate (limiting max download rate) is  $u_s + \sum u_i$



*time to distribute  $F$   
to  $N$  clients using  
P2P approach*


$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$




# P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks


**tracker:** tracks peers participating in torrent

**torrent:** group of peers exchanging chunks of a file



# P2P file distribution: BitTorrent

- peer joining torrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn**: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent





# BitTorrent: requesting, sending file chunks

## *requesting chunks:*


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

## *sending chunks: tit-for-tat*

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



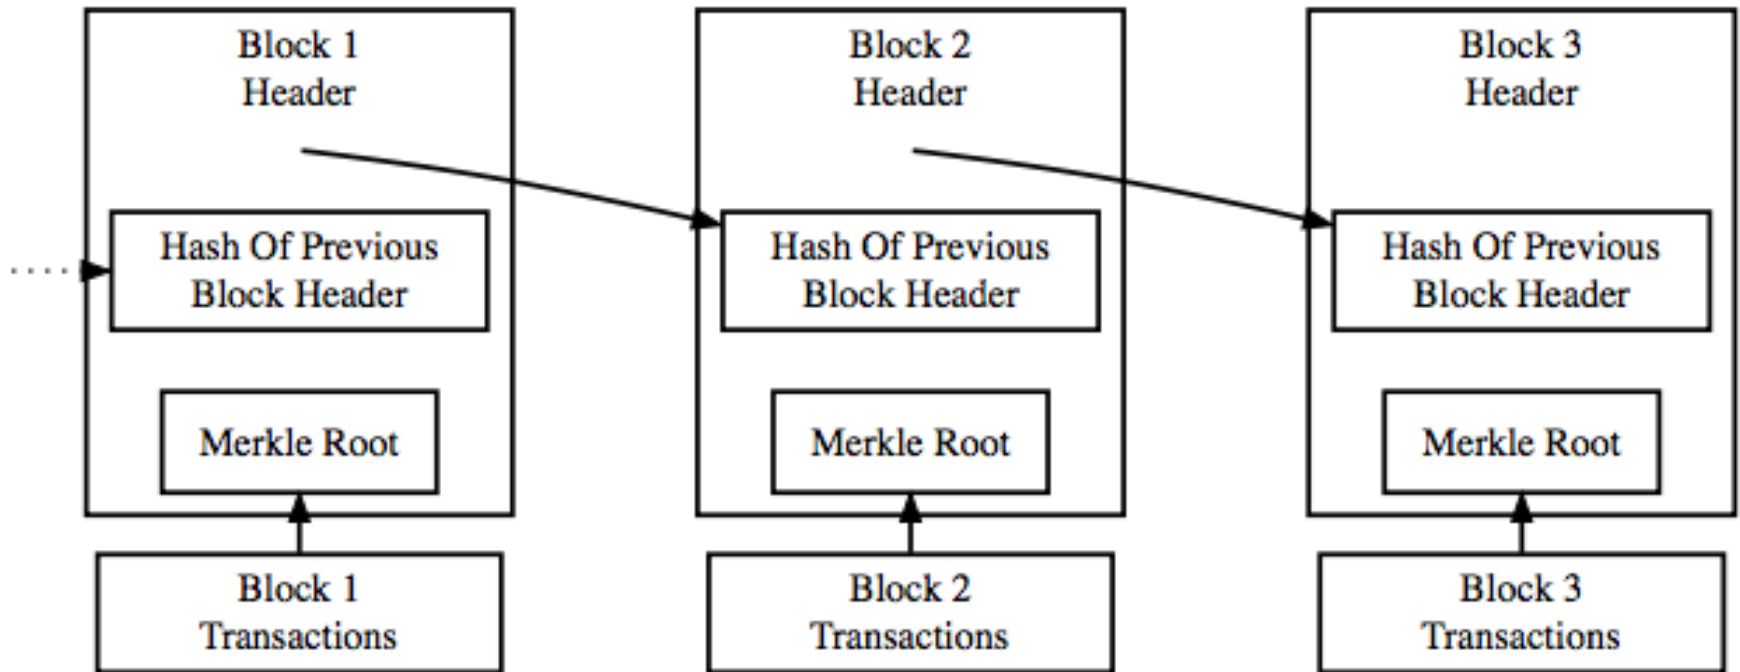
# Blockchain

- String of records, or blocks
- Each block contains a hash of the previous block in the chain
- Distributed across many nodes, fairly difficult to modify
  - Nodes are incentivized to reach consensus

# Blockchain continued

- Inherently decentralized
  - All nodes have a copy of the blockchain
  - Makes tampering more difficult, must use distributed consensus
- Proof of Work
  - Mining! Finding blocks is hard, verification is easy!
  - Difficulty in mining increases over time

# Blockchain illustrated



Simplified Bitcoin Block Chain

# Bitcoin

- First widely-adopted cryptocurrency
  - See <https://bitcoin.org/en/developer-guide>
- Built on blockchain transaction ledger
  - Contents within distributed database
  - Bitcoins are awarded for updating/verifying ledger
- Bitcoin are mined
  - Essentially, rewards for perpetuating the network



# Kademlia

## Network Programming

# DHT?

- Why is Kademlia called a Distributed Hash Table?
  - Partition the key space by using node IDs
  - Each ID stores a subset of the table (hence distributed)
  - How do we partition? We use XOR of ID and key as a hash
- So do nodes store <key,value> pairs as a hash table?
  - They can, but don't need to. Often use linked lists.



# Understanding Kademlia

- For HW3 we're not making an exact implementation but...
  - [details in next week's HW3 handout]
  - Will use Python / gRPC to do our RPCs
- Real Kad
  - UDP
  - Viewed as RPCs

# Understanding Kademlia

- For HW3 we're not making an exact implementation but...
- Every node generates an ID of length N
  - Fun fact - the behavior for ID collisions is undefined!
- To compute **distance** between IDs, simply use the bitwise XOR.
  - In Python/C/C++ that's operator<sup>^</sup>
- Keys are also length N. So we can compute node ID vs key **distance** the same way.

# Node ID Generation

- In real Kad, just pick from  $2^{160}$  randomly
  - If there is a collision, nodes are probably too far to see each other (unless  $k$  is large)

# OpenSSL SHA-1 Code (C, 1/2)

```
#include <openssl/evp.h>
EVP_MD_CTX *mdctx;
const EVP_MD *md;
char input[] = "12345";
unsigned char id[EVP_MAX_MD_SIZE];
int md_len;

//Create a context
mdctx = EVP_MD_CTX_new();
```

# OpenSSL SHA-1 Code (C 2/2)

```
/*Initialize our digest to use the md TYPE object
given by EVP_sha1 */
/*Need to do this every time we want to hash
something */
EVP_DigestInit_ex(mdctx, EVP_sha1(), NULL);
//Add to the hash
EVP_DigestUpdate(mdctx, input, strlen(input));
//Finalize the hash, store in md_value
EVP_DigestFinal_ex(mdctx, id, &md_len);

//Free the context
EVP_MD_CTX_free(mdctx);
```

# OpenSSL SHA-1 Code (Python 3)

```
#!/usr/bin/env python3
import hashlib
m = hashlib.sha1()
m.update(b"This is a str")
m.update(b"ring in two parts")
m.digest() #Gives the actual hash
```

# k and k-buckets

- Routing is split up into k-buckets. There are **N** of them.
- Each bucket  $0 \leq i < N$  holds nodes with  $2^i \leq \text{distance} < 2^{i+1}$
- If a node is in bucket  $i$ , then bit  $i$  is the first bit that was different (lowest bit on the left)
- For small  $i$ , there are very few possible values, probably an empty bucket.
- For large  $i$ , there are many possible values. Don't want to store them all.
- Only store up to  $k$  values

# k and k-buckets

- How do we pick which ones to kick out?
- Least Recently Used (LRU) list
  - If we follow the [Kad paper](#) the most recently used is the tail and the least recently used is the head
  - Head or tail is fine, as long as it's an LRU list
  - If you receive a message from a node, send it to the "most recent" end of the list
  - We don't want to update LRU when we send to a node, because it might turn out that node is dead. UDP means we won't get a "connection refused"




# k and k-buckets

## ■ Replacement Cache


- We won't do this in HW3
- In real Kad, can keep a number of nodes as "backup" - if a spot opens up in k-bucket, use a backup node
- Instead, only kick out if node fails to respond to a PING
- We won't even implement PING, real systems do though it may be very infrequent, e.g. once per 24 hours

# Visualizing k-buckets



Source: [https://commons.wikimedia.org/wiki/File:Dht\\_example\\_SVG.svg](https://commons.wikimedia.org/wiki/File:Dht_example_SVG.svg)

# Visualizing k-buckets



Source: <https://pdos.csail.mit.edu/~petar/papers/maymounkov-kademlia-lncs.pdf>

# Kademlia Protocol

## ■ Four RPCs

### ■ FIND\_NODE(id)

- Returns (IP, UDP Port, Node ID) for k closest nodes to id that are known by the remote node

### ■ FIND\_VALUE(key)

- Same thing as FIND\_NODE but if the node is storing a value for that key already, then it just returns the value

### ■ PING()

- Check if a node is responding

### ■ STORE(key,value)

- Store the key,value pair. Initiator asks k closest nodes to key.

# Doing Lookups

- “Concurrency” parameter  $\alpha$
- Pick up to  $\alpha$  nodes from closest non-empty k-bucket (we may only know fewer than  $\alpha$ )
- In parallel, make the same RPC (FIND\_NODE or FIND\_VALUE) to each of those nodes
- RPCs will return nodes, update our buckets and repeat the search with another  $\alpha$  nodes (but do not ask nodes we already asked)
- Stop when all k closest nodes have been asked

# Kademlia Performance

- Hopefully we can do a search in  $O(\log N)$  time, proving this is a little tricky
- $\alpha=1$  behaves similar to Chord for searching performance
- If node IDs are very far away, number of nodes is low, and  $N$  and  $k$  are not chosen well, it is possible not all nodes can reach each other.

# Question

- For  $N=5$ ,  $k=8$ , and a network with 32 nodes connected, which node IDs could be in bucket  $i=3$  for the node with ID = 10110?

You can put your answers in binary, hex, or decimal. (Hint: there will be 8 IDs)